

# Single Policy Multi Objective Q-Learning as applied to the Chinese Postman Problem.

Austin Byrd

## Abstract

The application of reinforcement learning to real world problem statements has become increasingly prevalent in artificial intelligence. As analytical methods break down on large scale systems, the solution space grows combinatorially. Recent breakthroughs in combinatorial optimization problems, such as Vehicle routing problems (VRP) (Nazari et al., 2018), give promising applications to reinforcement learning to problems that as of now require non-polynomial time algorithms to solve or approximate. The application has been applied to the Traveling Salesman Problem (TSP), which is a nodal routing problem, while less attention has been paid to edge and arc routing problems, such as the Chinese Postman Problem (CPP) until recently. Recent advancements have been made on CPP with extension of Load dependent costs CPP-LD by applying Graph-Attention based Deep Reinforcement Learning. The purpose of this paper is to use reinforcement learning to create an agent that can be embedded into a CPP valid graph and can learn on the graph to approximate an optimal policy to solve the CPP for that graph by using a Multi-Objective Q-learning framework.

## Introduction

The Chinese Postman Problem was first posed by a chinese mathematician named Kwan Mei-Ko in 1962 (Sokmen et. al, 2019). It encapsulates the issue of a Postman who must deliver the area's mail. To do this the Postman must start at the postoffice, go down every street, and after having gone down every street return to the postoffice.

This problem has many real world applications such as circuit testing, robotic path planning, inspections, distributed ledgers, etc.

The problem can be stated more formally as such:

$$G = (V, E)$$

Where  $G$  is a undirected weighted graph made of the set  $V$  of vertices, and the set  $E$  of weighted edges.

The objective of the CPP is to find the **shortest or lowest weight** path that

---

Copyright © 2022, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

1. Begins and ends at the same node ( $V_0 = V_f$ )
2. Crosses every edge in the graph (Path crosses all  $e \in E$ )
3. Minimizes the cost of the weighted edges taken in the walk (circuit).

The problem takes place in a graph, and the agent must begin at the start node and cross edges until all edges have been crossed, this objective is must be completed first before the agent returns to the start node to end the episode. The agent may cross edges multiple times in its path, but to complete a solution it must return to the start node after crossing all edges.

The problems main focus is upon Non-Eulerian graphs as it can cross edges multiple times. This is an important difference between the CPP and TSP.

The statement of the problem includes the optimality condition of **shortest or lowest weight** path.

Analytical methods have approximately solved the problem and heuristic methods have improved upon the most optimal attainable solution space (Pan et al., 2024). The continuing issue with those methods is that as the problem graph grows the algorithm runtime grows exponentially. This poses the question, can a reinforcement learning algorithm match the approximal optimality of analytical methods when applied to a traditional Chinese Postman Problem? And, if so, can it be altered to generalize to any CPP valid problem graph?

## Background

Reinforcement learning is a type of machine learning that relies upon the agent or machine to learn from past actions and rounds of actions to create a policy that can guide its future actions. This allows for an agent to be immersed into an environment and make actions. From those actions values can be calculated for each action taken from a given state in the environment. This ability to learn from the environment by being immersed in the environment and making actions allows for knowledge to be collected and applied when the same choices are made again.

---

**Algorithm 1: Chinese Postman Problem Algorithm**

---

**Input:**  $G = V, E$  where  $G$  is a Non-Eulerian Graph**Output:** Eulerian Graph

- 1: Find the degree for each node in the graph.
  - 2: Pair odd degree nodes together greedily minimizing path distance.
  - 3: Add edge between each pair of odd degree nodes. Making every node degree even.
  - 4: Find Eulerian Circuit for the graph.
  - 5: Sum the length of the Eulerian circuit.
  - 6: Apply Hierholzer's Algorithm
  - 7: Return Eulerian Circuit length.
- 

These tools allow for the formation of a policy that can outperform analytical methods on certain problems.

The application of reinforcement learning to arc routing problems could yield policies that human ingenuity could not make, but only could be made by large amounts of experience.

## Related Work

The Chinese Postman Problem is a NP-Hard problem that has been attempted with analytical approximations and heuristically (Sokmen, et. al. 2019). Due to the NP-Hard nature of the question the time complexity of algorithmic solutions grow exponentially with the number of odd degree edges in the graph. See Algorithm 1. for the traditional algorithm.

The time complexity of the method outlined in Algorithm 1 is set by the basic operation that is step 3 of the algorithm as the number of possible pairs of odd degree nodes grows combinatorially. This has been approximated via greedy methods but this does not provide exact solutions.

Real world application of the Chinese Postman Problem include Load dependent costs (CPP-LDC). CPP-LDC relies upon that the load to be delivered gradually diminishes as the edges are crossed, lowering the cost of edges that are traversed later in the delivery cycle. This can have resource savings in real applications of fuel usage and travel time.

These real world applications have inspired many methods of solving the CPP and many more complicated problem spin offs of the CPP. These more expansive methods have been formulated for more complicated versions of the CPP such as CPP-LD, Rural Postman Problem, Windy CPP, Directed CPP, Time Dependent CPP, Multiple CPP, etc. (Sokmen et al., 2019) (Christofides et al., 1981) (Pan et al., 2024) (Corberan et al., 2018) (Corberan et al., 2021) (Tran et al., 2023) (Terefe, 2015). These problems have many different methods that are used to approximate optimal solutions. A variety of methods are included in a few notable libraries <sup>1</sup> (Corberan et al., 2021).

<sup>1</sup>(<https://github.com/Olibear/ArcRoutingLibrary>  
<https://www.uv.es/corberan/instancias.html>) and

The huge majority of these methods have been analytical, and recently heuristic and meta-heuristic methods have great promise including reinforcement learning.

The application of a model known as Arc-DRL has been the most effective (Tran et al., 2023). Arc-DRL is an Arc based Graph traversal algorithm that uses Graph-Attention based Deep Reinforcement Learning to train a Graph-Attention Encoder that produces embeddings to pass through some number (N) of 2 layer attention layers and then to the decoder. The decoder is intended to reconstruct the optimal tour of the input graph. This model has done exceptionally well at solving the CPP-LD. The neural network was trained using 2 datasets from a related experiment (Corberan et al., 2018).

Additional models that have been used on this problem include: the Genetic algorithm, Ant Colony algorithm, Greedy Heuristic Construction, Iterated Local Search, and Variable Neighborhood Search as well as others.

Many of these models utilize cutting edge technologies and models in the field of Machine Learning and some a multi-model hybrid algorithm that have many layers of neural network (Vaswani et al., 2017).

The application of a concept known as the Pareto Front has grown in popularity for problems such as the CPP. The Pareto Front is a subset of the solutions come to by an algorithm that are most optimal for balancing multiple objectives (Van Moffaert, Nowe 2014). This most optimal subset of solutions is saved and only superceded when solutions of higher quality and optimality for both objectives are observed. Then the Pareto Front is updated to the most optimal solution and the algorithm continues searching.

The CPP being a two stage problem requires a balance of two objectives. The first being to cross all the edges of the graph, and the second being to return to the start node. Therefore there are two objectives that must be accomplished. This multi-objective problem requires Multi Objective Markov Decision Process (MOMDP) and Multi Objective Q-learning (MOQ) to solve (Yang, et. al 2019) (Hayes, et. al 2021). MOQ and MOMDP have been formulated and applied to many problems in Reinforcement Learning, but until now MOQ has rarely been applied to the CPP(Majumder 2019).

## Methodology

The type of reinforcement learning algorithm proposed for implementation is a Multi-Objective Q-Learning model to learn on a specific graph. This method is implemented by transforming the CPP into an MOMDP. To transform the Chinese Postman Problem into a Multi-Objective Markov Decision Process there must be a clear set of States, Actions, Vectorized Rewards, a Scalarized Return function, a discount factor, a learning rate, and a policy (Van Moffaert, Nowe 2014). Python was used for program development.

## Multi Objective Reinforcement Learning

MOQ learning is different from traditional Q-learning in many ways. The most important is the reward vectorization,

the Q value scalarization, and the alteration of the action selection method. Additionally the transition function is required.

These will be applied to the Chinese Postman Problem via:

## Objectives

The agent has two objectives the first objective  $\Xi$  is to cross all edges in the graph.

The second  $\Upsilon$  is to return to the start state.

$$O = \{\Xi, \Upsilon\}$$

## States

Set of states  $S$ : Nodal Position of Agent in the graph  $v \in V$

Goal State  $s_f$ : When all other nodes are reached and the agent has returned to the initial state.  $s_0 = s_f$

## Actions

Set of actions  $A$ : The agent can accross any Edge connected to current Node Position  $e \in E_s$

Actions in MOQ are made via a Transition function  $T$ .  $T$  is a probability distribution function among the possible actions in a state, where a state, action pair is input and the output is the probability of that action being taken from that state following the Q values  $Q_\pi^O(s, a)$  over all objectives.

## Rewards

Set of rewards  $R = \{100, 0, -5\}$  Agent receives a two rewards, one for each objective  $o \in O$ .

For objective  $\Xi$  the agent recieves a reward of 100 for reaching the start/goal state, and a reward of -5 for any other states.

For objective  $\Upsilon$  the agent recieves a reward of -5 for crossing any edges already crossed. The agent recieves a reward of 0 for crossing edges that have not been crossed yet. Lastly, the agent recieves a reward of 100 for crossing all the edges in the graph.

This has room for manipulation to tune the model further.

## Discount Factor

A discount factor of  $\gamma = .8$  will be applied, a low discount factor is used to highly encourage the agent to immediately accomplish the goal of solving the CPP.

## Learning Rate

A learning rate of  $\alpha = .9$  is applied and gradually decreased until it is at .1 to start learning quickly and rapidly slow the learning.

## Return Function

Return Function  $Q_\pi^O(s, a)$  etc. (For a given Objective there are distinct functions for the Value, Goodness, or Quality Expectation of a state or given action in a state)

For MOQ with two objectives there are two Q tables one for each objectives Q values. Both Q tables are kept in one larger vectorized Q table. Each objectives Q tables are

trained on that objectives return. This updates the objectives Q tables simultaneously on a single state, action pair but updates both Q tables on that action depending on the return for that action given the objective.

The return functions are trained on the episodes of training the agent goes through. The functions are trained on the averaged value of a state in an episode by utilizing the discount factor to reduce the reward of a state if it takes many actions to reach the goal state rather than it taking fewer actions to reach the goal state via a different action/state. This combined with small negative rewards highly incentivize the agent to end the episode quickly for highest reward. This numerically ranks the states in a vectorized mapping between states and values for given objectives.

$Q_\pi^O(s, a)$  is a matrix of values mapping state action pairs to an expected return value **for an objective**.

## Policy

Policy  $\pi$ : Initially equally distributed random actions are selected at any given state. This initial random policy is used to collect the data needed to calculate  $Q_\pi^O(s, a)$  which is then updated as the quality of state action pairs are calculated for both objectives.

The combination of these objective Q tables into one final policy can be done in multiple ways. The chosen method is via Pareto Dominance, a Nash Equilibrium Dominant Strategy analagous policy selection method. That is to identify the most optimal solutions of the Pareto Front.

## Methodological Philosophy

The methodological outlook is to formulate a MOQ-learning RL model. Test the models capabilities on a curated CPP with a known optimal solution. Then to simulate the MDP. This will train the return function on episodes of training from a state space; meanwhile collecting a Pareto Front. This initial training of the model on one graph that has a known solution ensures the evaluation of the model's capabilities before being applied to larger CPP instances.

Once the algorithm is successful upon examples of the CPP, the model will be applied to create policies upon more complex graphs with larger number of nodes and edges for the algorithm to run upon. This creates a different model for every graph.

The final goal will then be to attempt to create a Multi-Policy algorithm using MOQ-learning that will create policies to improve the optimality of the Pareto Dominant strategy found by the algorithm.

## Evaluation

The evaluation of the reinforcement learning algorithm will be carried out by comparing the solutions of the trained final policy; that was formulated from Q values informed by the Pareto Front, to a solution calculated with the traditional analytical approximation methods for solving CPP.

The evalution method utilized is a comparison made between Single Policy MORL, and the traditional analytical algorithm.

- The benchmark of Single Policy MORL methods will be compared to the results quantified by finding the performance ratio between analytically solution methods and each of the MORL methods.
- The goal is for the MORL algorithms policies performance ratio to be  $\rho \leq 1.1$  or within 10% of the approximations from Algorithm 1 respectively.

## Results

The evaluation of the model is done with 2 graphs.

### Graphs

1 curated small graph of 5 nodes and 8 edges

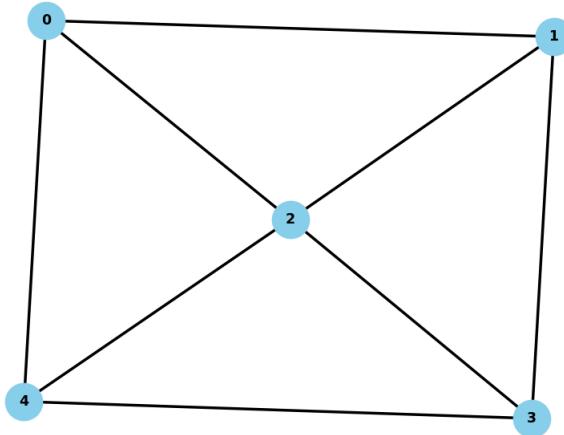


Figure 1: Small Graph

1 randomly generated large graph of 25 nodes and 40 edges

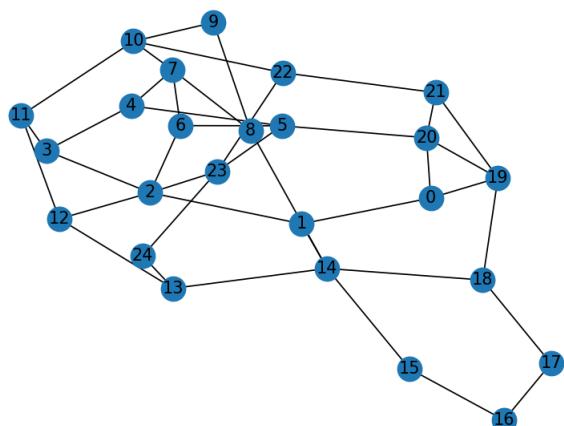


Figure 2: Large Graph

The solution space that is observed in the final policies and Pareto front of single policy MORL agents solving the CPP are approximately equivalent to traditional analytical methods. The training of the agents upon a single graph to solve for a solution The Pareto Frontier is a subset of soluitons that are superior solutions observed from all episodes of training. The Pareto Dominant solutions observed in the sample of 1000 episodes shows that the models continuous improvement converges at the optimal solution as the number of episodes approaches infinity. Furthur episodes have the chance to further improve on the solution space as a large sample of the space allows for better assurance of optimality.

### Pareto Frontier

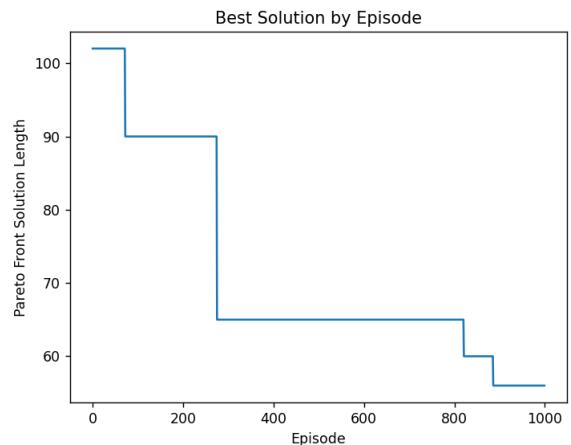


Figure 3: Pareto Front Best Policy by Episode Large Graph, optimal = 49

$$\rho = 52/49 = 1.06$$

The Pareto Dominant solution that is concluded at the end of the 1000 episodes for the large graph is 52, which is 3 actions away from the analytical solution of 49 actions required to complete the CPP of the large graph.

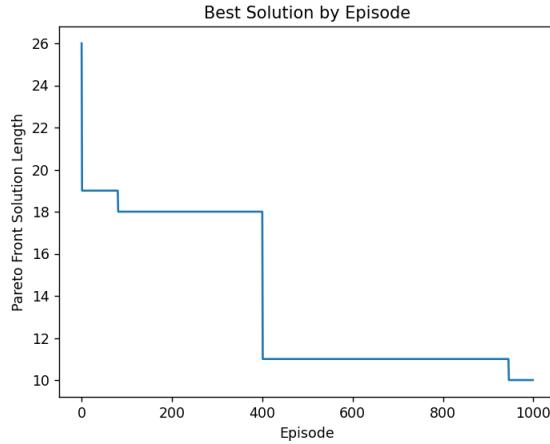


Figure 4: Pareto Front Best Policy by Episode Small Graph, optimal = 10

$$\rho = 10/10 = 1$$

The small graph's final Pareto Dominant solution is 10 which is equivalent to the solution derived from the analytical method. This shows that the algorithm finds the optimal solution for small graphs, and indicates that the algorithm when applied to larger problem graphs will converge to at least the approximately optimal solution guaranteed by analytical methods.

## Cummulative Reward

The cummulative reward evaluation of a Multi-Objective RL model is slightly altered from the traditional RL cummulative reward analysis. This is exacerbated by the Single policy method used which creates a greedy combination of the policies formed for each objective individually.

This is exhibited in the models cummulative reward as a very slow growth with a large amount of variance in the objective reward early on as the agent uses the initial epsilon greedy method to explore. The agent explores (ignores one objective in favor of the other) causing volatility in either objectives reward, and as the training continues the reward settles as the agent shift to a more predominantly exploitative action selection policy.

This requires the agent to synthesize a Pareto Dominant policy out from the maximum position in between the rewards of the two objectives. This combination of the policies in a greedy way leads to final policies that allow for the consideration of two objectives at the cost of focusing upon achieving a single objective.

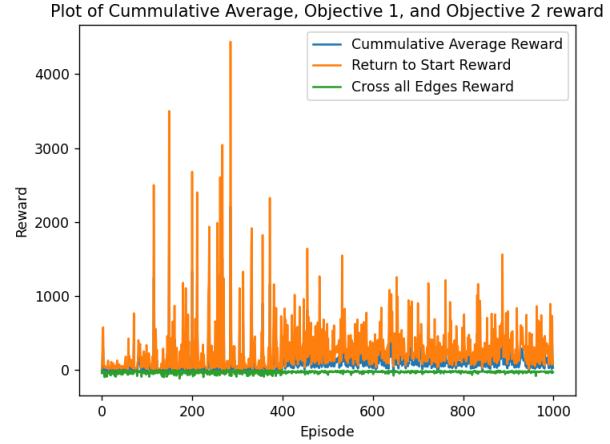


Figure 5: Reward by Objective and Average Large Graph

The cummulative reward grows as the training progresses. This indicates that the single policy is being improved by training alongside the Pareto Frontier solutions as well as that the Pareto Frontier is expanded with exploration of the solution space. The stability of the solution is additionally stabilized as the variance of the two objective rewards is calming by the end of the training. The settling of the values are showing that as the training progresses the agent explores less and instead exploits.

The slow progress of the training is due to the volatility of early reward to the Return to start objective.

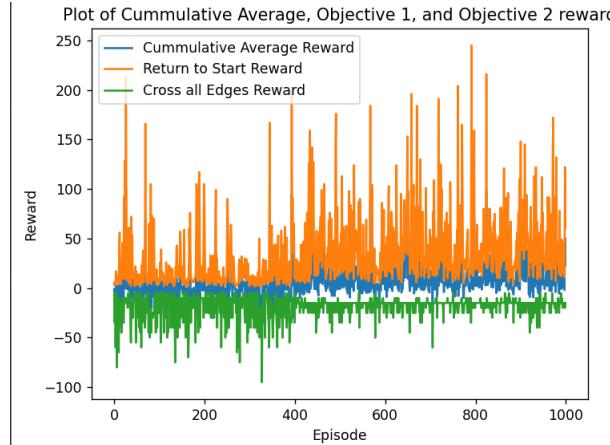


Figure 6: Reward by Objective and Average Small Graph

The cummulative reward of the agent over training upon the small graph trends to a higher reward. This indicates the agent is learning from the environment. The final policy being a greedy combination of the two objectives and learning from the Pareto frontier does not ensure that the final policy is optimal only that it was informed by the Pareto frontier. This exemplifies the way that the final policy must be formulated from two contrary policies. This still allows the agent to formulate policies that approximate analytically derived approximate solutions. The bootstrapping nature of this al-

gorithm is slowed by the balancing of the two objectives.

Single Policy MORL values learn to balance between the two objective rewards to find a golden mean policy between the two objectives separate optimal policies. As the agent learns, the volatility of average reward becomes higher as the Q values and exploitation rise ensuring the agent balances between the two objectives. This forms the single policy MORL

The single policy MOQ is the best solution that incorporates both objectives. The optimal solutions for the single policy is the Pareto Dominant solution of the Pareto Frontier. The Pareto Frontier is saved as the best solutions that have been observed in all episodes of training. The Pareto Dominant solution is the solution of an observed episode of training that achieved both objectives in the lowest number of actions.

## Summary of Results

The Chinese Postman Problem as approached by Multi-Objective Q-Learning Reinforcement Learning yields solutions to Chinese Postman Problems that are within the +/- 10% approximation goal of evaluation initially set forth. This accomplished the initial project goals, but did not exceed the goals as the runtime to solve the CPP using RL is an order of magnitude slower than using the traditional non-polynomial solution. This result indicates improvement must be made in the solutions for the application of MORL to larger instances of CPP to be a worthwhile endeavor for expanding the solution space of CPP. Preliminary results indicate potential in further investigation as these early results resemble analytical methods and further work could improve even further upon the results.

## Possible Improvements

This model could be improved and expanded in many ways. One aspect of the traditional Q-learning algorithm is that it overestimates the values (Van Hasselt, Guez, and Silver 2016). It is shown that by using alterations of the traditional algorithm; such as Double Q-learning, or Self-Correcting Q-learning that the learned values can more closely estimate the values in the Q table. Applying these methods to Multi-Objective Q learning could have useful applications. This could be implemented by using the advanced Q methods on each of objectives Q tables within  $Q_\pi^O(s, a)$ . Another improvement to this model would be to add weights to the edges and attempt to solve the problem when the weights of each edge are not all uniformly 1. This would add additional real world application to the model as not all roads are of the same length, steepness, or quality of passage. This inequality of edge cost would additionally be more faithful to the original problem statement of CPP.

## Conclusions

The research conducted in this experiment has been focused upon using Multi-Objective Single Policy Q-learning for application to the Chinese Postman Problem. The objectives of the agent are to cross all edges, and then to return to the start node. The application of these two objectives to

an agent solving the CPP and then to formulate a single policy from the Q values trained on many episodes and informed by Pareto front solutions yields approximate solutions that match analytical methods for solving the CPP. Future explorations have been proposed for expanding the solutions to Multi-Policy MOQ as well as weighted graphs. This would allow for further elucidation of the efficacy of Multi-Objective RL to the CPP.

## Timeline

- 3/31 Complete creation of Multi Policy Q learning algorithm
- 4/14 Complete training of Expand all models to weighted edges problem
- 4/28 Complete Evaluation of Models with weighted edges
- 5/1 Final Presentation

## References

Tran, Cong Dao, and Truong Son Hy. 2023. Graph Attention-based Deep Reinforcement Learning for solving the Chinese Postman Problem with Load-dependent costs. arXiv preprint arXiv:2310.15516

Sokmen, Ozlem and Emec, Seyma and Yilmaz, Mustafa and Akkaya, Gokay. 2019. *An Overview of Chinese Postman Problem.*, International Conference of Advanced Engineering Technology

Pan, P., Zhu, H. 2024. *Approximation algorithms for the restricted k-Chinese postman problems with penalties.* Optim Lett 18, 307–318 . <https://doi.org/10.1007/s11590-023-01992-z>

Farahani, Reza Zanjirani, 2012. *Graph Theory for Operations Research and Management: Applications in Industrial Engineering: Applications in Industrial Engineering.* IGI Global.

Christofides, N., Campos, V., Corberán, A., and Mota, E. (1981). *An algorithm for the rural postman problem.* Report, International Conference for Operations Research.

Corberán, Á., Eglese, R., Hasle, G., Plana, I., and Sanchis, J. M., 2021. *Arc routing problems: A review of the past, present, and future.* Networks an International Journal.

Corberan, A., Erdogan, G., Laporte, G., Plana, I., and San- chis, J. (2018). *The chinese postman problem with load- dependent costs.* Transportation Science

Terefe, Abel, 2015. *Arc routing in household waste collection,* Aalto University

Van Hasselt, H.; Guez, A.; and Silver, D. 2016. *Deep reinforcement learning with double q-learning.* In Proceedings of the AAAI conference on artificial intelligence, volume 30.

Hayes, C.F., Rădulescu, R., Bargiacchi, E. et al. 2021. *A practical guide to multi-objective reinforcement learning and planning*. Auton Agent Multi-Agent Systems

Yang, R., Sun, X., Narasimhan, K. 2019. *A generalized algorithm for multi-objective reinforcement learning and policy adaptation*. Advances in neural information processing systems

Van Moffaert, Nowé, A. 2014. *Multi-objective reinforcement learning using sets of pareto dominating policies*. The Journal of Machine Learning Research

Gábor, Z., Kalmár, Z., Szepesvári, C. 1998. *Multi-criteria reinforcement learning*. ICML

Natarajan, S., Tadepalli, P. 2005. *Dynamic preferences in multi-criteria reinforcement learning*. In Proceedings of the 22nd international conference on Machine learning

Majumder, S., Kar, S., Pal, T. 2019. *Uncertain multi-objective Chinese postman problem*. Soft Computing

### **Acknowledgments**

We would like to thank Tom Lehrer for all of his support.  
Thank you for reading my paper