


Starting where Lecture 3 left off

Dynamics of an RL Problem

① State-action-state or transition table

From state

	actions				
	↑	↓	←	→	↻
s_1	NA	s_4	NA	s_2	NA
s_2					
s_3					
s_4					
s_5					
s_6					



② State-action-immediate reward table or Reward table

From state

	action				
	↑	↓	←	→	↻
s_1	NA	0	NA	0	NA
s_2	NA	0	0	100	NA
s_3	NA	NA	NA	NA	0
s_4	0	NA	NA	0	NA
s_5	0	NA	0	0	NA
s_6	100	NA	0	NA	NA


Non-determinism (CS4700)

Actions

→ : go left action

instead use

Try → : Try to go left action



or could use 100% → s_2 and 0%, 0%

Reward Non-determinism

For the same action at the same state the rewards are stochastic


Most complex non-deterministic case

a combination of these two

Goal of RL

To learn a better policy, possibly the ideal/best

There may be many or infinite possibilities but "one" best policy may be



What is the best action to perform in each state

Policy π


initial policy π_0

$$\pi_0 = p(a|s) \forall s \in S$$

assume an equi-possible policy [or it may be random]
all actions possible in a state are equally probable
can be written as a table in the simple case

state

	action				
	↑	↓	←	→	↻
s_1	0.5	0	0	0.5	0
s_2					
s_3					
s_4					
s_5					
s_6					



to perform in each state

Policy may include offline knowledge


Learned "best" policy

state	action
s_1	→
s_2	→
$\vdots$	
s_6	

Reinforcement learning

given the dynamics of the problem and T_0 (random, e.g. vi-prob, expert-based) learn a "best" policy.

given the dynamics of the "maze" use e.g. vi-prob policy (say)



need to find out how good state s_t is for the RL agent wants to maximize the cumulative immediate rewards

Return G_t of state s_t

$$G_t = R_{t+1} + R_{t+2} + \dots + R_T \quad \left[\begin{array}{l} \text{sum of immediate rewards} \\ \text{simple cumulative reward} \end{array} \right]$$

$G_{1,1}$
↑
episode

for this maze all $G_t = 100$

would need to track steps to make this learn

could use this:

$$\text{Return} \Rightarrow G_t = R_{t+1} + \underset{\substack{\uparrow \\ \text{discount factor} \\ \text{gamma}}}{\gamma} R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t} R_T$$

where $\gamma = 1 \geq \gamma \geq 0$

assume $\gamma = 0.9$

$$G_1 = 0 + \gamma * 100 = 0 + 90 = 90$$

$$G_2 = 0 + \gamma * 0 + \gamma^2 * 0 + \gamma^3 * 0 + \gamma^4 * 100 = 65.61$$

$$G_3 = 0.002656$$

Value of a state

$$V_{\text{episode}_i}(x_i) = G_i \text{ episode}_i$$

In general $V(s_t) = G_t$ in some episode

What is the generic value of a state in an episodic RL problem?

$$V_{\pi}(x) = \underset{\substack{\uparrow \\ \text{Policy}}}{\mathbb{E}} \left[\underset{\substack{\uparrow \\ \text{mean value}}}{G_t} \mid s_t = x \right] \times \mathbb{P}(s) \Rightarrow \frac{1}{\# \text{ episodes}} \left[R_{t+1} + \gamma R_{t+2} + \dots \right]$$

Naive algo to learn $V_{\pi}(s)$

For every state s $\begin{cases} \text{no_visits}[s] = 0 \\ V(s) = 0 \end{cases}$

for $i = 1$ to no_episodes

start episode_i from a start state

finish episode_i

for $t = 1$ to T_i

$$G_{t,i} = R_{t+1,i} + R_{t+2,i} + \dots$$

$$V(s_t) = \frac{\text{no_visits} * V(s_t) + G_{t,i}}{\text{no_visits} + 1}$$

Issues with naive algo

- in big environments, this is very very slow
- convergence may not happen
- deterministic
- may not visit all states